

Predicting stock prices with LSTM



Alexandre Xavier

Follow

Jan 22, 2019 · 8 min read



The Portuguese version of this article is available in [Prevendo valores de ações com LSTM](#)

In this article we will use Neural Network, specifically the **LSTM model**, to predict the behaviour of a **Time-series** data. The problem to be solved is the classic stock market prediction. All data used and code are available in [this GitHub repository](#).

Although this is indeed an old problem, it remains unsolved until today. The truth is quite simple: the stock price is determined by several factors, and its historical price is just a small fraction of it. As a result, predicting the price behaviour is a really difficult problem, unless you are a [monkey playing darts](#).

Summary

First, I will introduce the dataset with some **data visualisation**. Then, I will briefly discuss how **difficult it is** to predict the stock market behaviour by using the moving average algorithm and showing its limitations. Next, a short introduction to the concept of Recurrent Neural Networks and LSTM, followed by a LSTM example of **predicting the stock price for a single company**. Finally, I will show the LSTM used to **predict the price of four companies** at the same time and compare the results to see if the prediction improves as we use more companies at the same time.

Data visualisation

The dataset was downloaded from [Yahoo Finance](#) in CSV. It has the stock price of four companies in the period between 01/08/2010 and 01/07/2019. We will refer to them as company A, B, C and D.

The basic step is to open the CSV file using Pandas. In a first look at the data, we have:

```
df_A = pd.read_csv('data/company_A.csv')
df_A['Date'] = pd.to_datetime(df_A['Date'])
df_A.tail()
```

	Date	Open	High	Low	Close	Adj Close	Volume
2259	2018-12-31	158.529999	159.360001	156.479996	157.740005	157.740005	35003500
2260	2019-01-02	154.889999	158.850006	154.229996	157.919998	157.919998	37039700
2261	2019-01-03	143.979996	145.720001	142.000000	142.190002	142.190002	91312200
2262	2019-01-04	144.529999	148.550003	143.800003	148.259995	148.259995	58607100
2263	2019-01-07	148.699997	148.830002	145.899994	147.929993	147.929993	54693100

```
plt.figure(figsize = (15,10))

plt.plot(df_A['Date'], df_A['Close'], label='Company A')
plt.plot(df_B['Date'], df_B['Close'], label='Company B')
plt.plot(df_C['Date'], df_C['Close'], label='Company C')
plt.plot(df_D['Date'], df_D['Close'], label='Company D')


plt.legend(loc='best')
plt.show()
```



Moving Average

A classic algorithm used in this problem is the Moving Average (MA). It consists in calculating the average of the m past observed days and use this result as the next day prediction. To demonstrate, here is an example of a moving average using m as 10 and 20 days on company's A closing price.

```
df['MA_window_10'] = df['Close'].rolling(10).mean().shift() #shift so
the day we want to predict won't be used
df['MA_window_20'] = df['Close'].rolling(20).mean().shift()
```




When we try to forecast the closing price 10 days ahead using Moving Average, the results are as follows:



Notice that each red line represents a 10 day prediction based on the 10 past days. For this reason, the red line is discontinuous.

Using a slightly fancier algorithm, the Exponential Moving Average (EMA), we achieve some small improvement:



Contrasting the MA and EMA:




This approach is simplistic. What we really want is to predict \hat{n} days ahead to see stock the future behave, and both MA and EMA fail in this task.

Recurrent Neural Network (RNN)

To understand an LSTM Network, we need to understand a Recurrent Neural Network first. This kind of network is used to recognise patterns when past results have influence on the present result. An example of RNN usage is the time-series functions, in which the data order is extremely important.


In this network architecture, the neuron uses as input not only the regular input (the previous layer output), but also its previous state.



It is important to notice that \hat{h} represents the neuron state. Therefore, when in state \hat{h}_{-1} , the neuron uses as input the parameter \hat{x}_{-1} and \hat{h}_{-0} (its previous state). The main problem of this model is the memory loss. The network older states are fast forgotten. In sequences where we need to remember beyond the immediate past, RNNs fail to remember.

LSTM Network


An LSTM Network has its origin in a RNN. But it can solve the memory loss by changing the neuron architecture.




The new neuron has 3 gates, each with a different goal. The gates are:

- *Input Gate*
- *Output Gate*
- *Forget Gate*

An LSTM Neuron still receives as input its previous state:





LSTM to predict a single company


At last, let's use an LSTM to predict the behaviour of company A alone.

But first, consider the following parameters. We want to predict the n days ahead (`forward_days`) having as input the m past observed days (`look_back`). So, if we have an input of m past days, the network output will be the prediction for the n next days. We will split the data in Train and Test. The test will be composed of k periods (`num_periods`), in which every period is a series of n days prediction.

```
look_back = 40
forward_days = 10
num_periods = 20
```

Now, we open the CSV file with Pandas and keep only the columns we will use, that are the date and the closing price. The initial closing price graphic for company A is:

```
plt.figure(figsize = (15,10))
plt.plot(df)
plt.show()
```



In sequence, we scale the input, split the data in Train/Validation and Test and format it to feed the model. All the detailed process can be found at my [GitHub](#).

Now, we build and train the model.

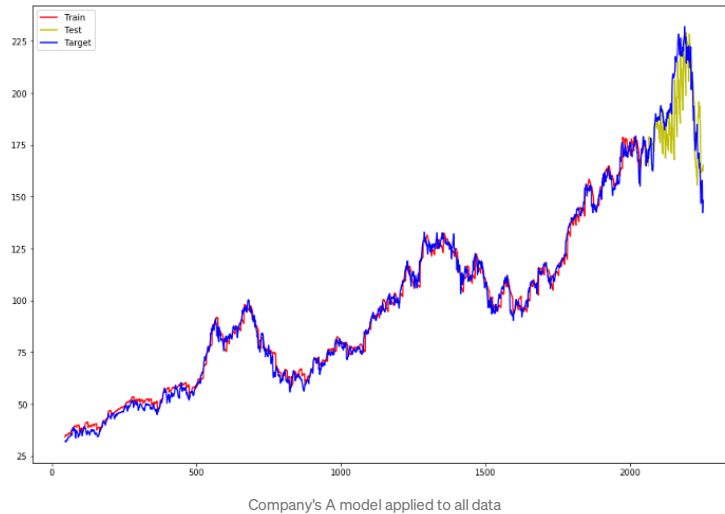
```
NUM_NEURONS_FirstLayer = 128
NUM_NEURONS_SecondLayer = 64
```

```
EPOCHS = 220
```

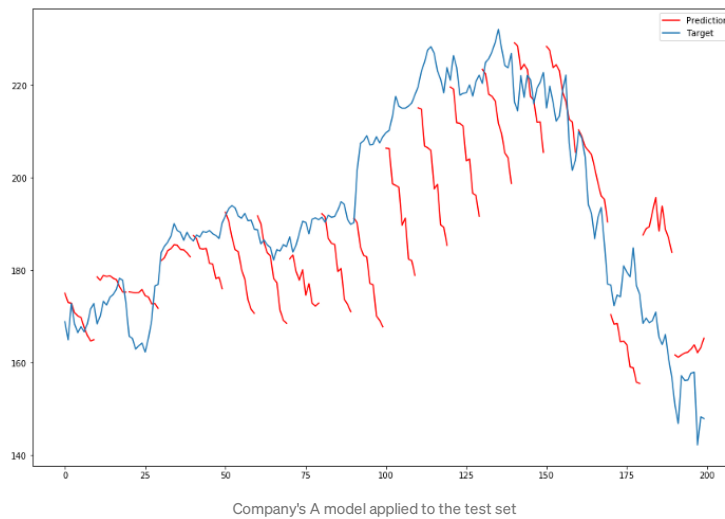
```
#Build the model
model = Sequential()
model.add(LSTM(NUM_NEURONS_FirstLayer,input_shape=(look_back,1),
return_sequences=True))
model.add(LSTM(NUM_NEURONS_SecondLayer,input_shape=
(NUM_NEURONS_FirstLayer,1)))
model.add(Dense(foward_days))
model.compile(loss='mean_squared_error', optimizer='adam')

history = model.fit(X_train,y_train,epochs=EPOCHS,validation_data=
(X_validate,y_validate),shuffle=True,batch_size=2, verbose=2)
```

The result obtained is:

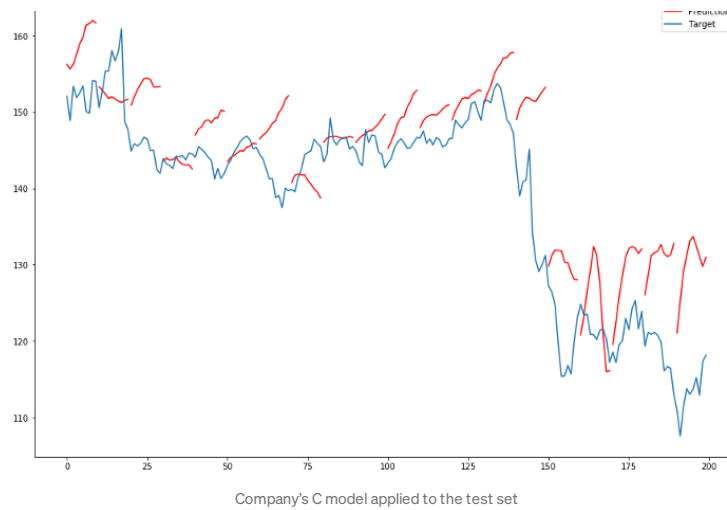


Looking the test set closer:



Notice that each red line represents a 10 days prediction (`forward_days`) based on the 40 past days (`look_back`). There are 20 red lines because we choose to test on 20 periods (`num_periods`). Thats why the prediction line in red is discontinuous.

By repeating the same process for all companies, the best result on the test set was the prediction for company C:



Even though this was the best model, the results are far from great. The reasons that may have caused this result are many. Some of them may be:

- Only the historical data of closing prices is not enough to predict the stock price behaviour
- The model could be improved

LSTM to predict four companies

Finally, we will use the LSTM model to predict the behaviour of all four companies together, A, B, C and D, and contrast with the single LSTM company results. The goal is to analyse if using the data from several companies can improve the individual prediction of every company.

It is important to point out that all four CSVs have the same dates. This way, the network won't be fed with future information from a company to predict the value of another company.

Initially, the data is:



After the data normalisation and formatting to feed the model, the model was trained:

```

NUM_NEURONS_FirstLayer = 100
NUM_NEURONS_SecondLayer = 50
EPOCHS = 200

#Build the model
model = Sequential()

model.add(LSTM(NUM_NEURONS_FirstLayer,input_shape=(
look_back,num_companies), return_sequences=True))
model.add(LSTM(NUM_NEURONS_SecondLayer,input_shape=(
NUM_NEURONS_FirstLayer,1)))
model.add(Dense(foward_days * num_companies))
model.compile(loss='mean_squared_error', optimizer='adam')

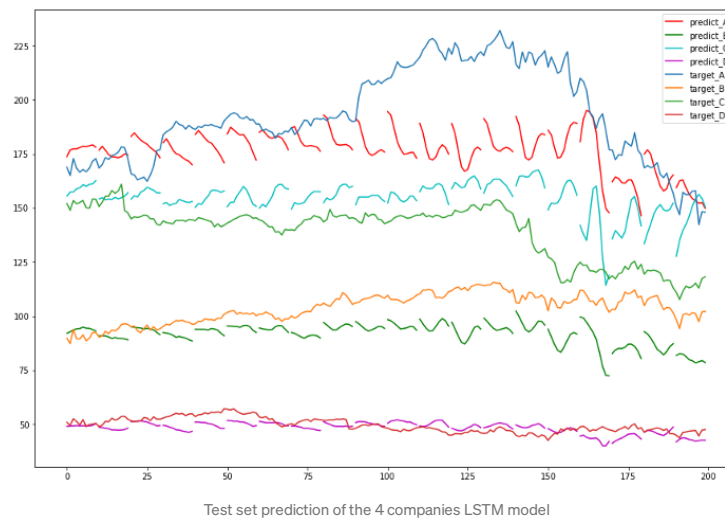
history = model.fit(X_train,y_train,epochs=EPOCHS,validation_data=(
X_validate,y_validate),shuffle=True,batch_size=1, verbose=2)

```

The result was:







Looking the test set closer:

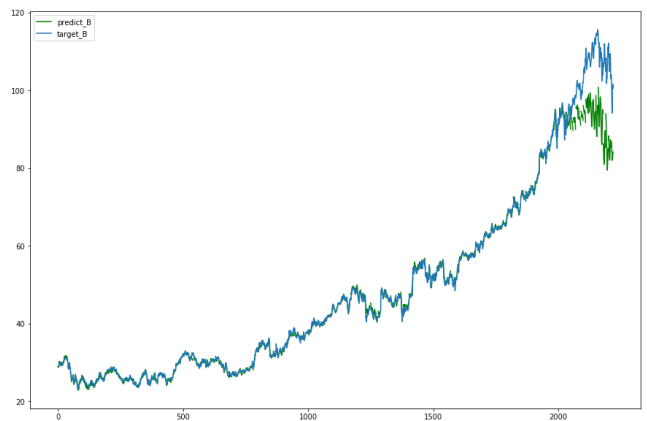




Time to contrast the results. The results of a the single company LSTM is shown on the left, and the result of the four company LSTM on the right. The first line shows the prediction in the test set, and the second on all the data set.




Company A



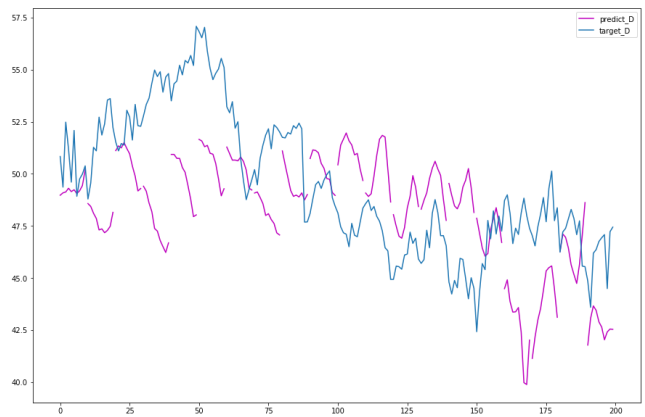

Company B



Company C



Company D



Conclusion

It is not possible to predict the stock market behaviour using only its historical price. The LSTM prediction is far from acceptable. Even when using the historical price of several companies, the prediction got worse.

References

<https://www.datacamp.com/community/tutorials/lstm-python-stock-market>

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

<https://towardsdatascience.com/train-validation-and-test-sets-72cb40cba9e7>

<https://machinelearningmastery.com/diagnose-overfitting-underfitting-lstm-models/>

<https://machinelearningmastery.com/reshape-input-data-long-short-term-memory-networks-keras/>

Thanks to .

Machine Learning

Lstm

Artificial Intelligence

Stock Market

👏 225

💬 3

🐦 in f ↗



WRITTEN BY
Alexandre Xavier

Follow



Neuronio

Neuronio is a Brazilian company that creates Deep Learning solutions and offers consulting services. At Medium, we write about machine learning and deep learning. #ai #deeplearning #machinelearning

Follow

More From Medium

**Technical Paper:
Enhancing a facial
recognition system via a
deep learning base
model as a...**

Fortune Okorji in Seamfix
Engineering



**Stochastic Gradient
Descent By Tensorflow
and Keras Framework**

Pattharadet Vachirapuchai



**"Body": using your self as
a search tool**

anselan in Random Studio



**Gradient Descent in
Machine Learning:
Simplified**

Nirmalya Misra in Python in
Plain English



**Does Data Science relate
to Machine Learning and
Deep Learning? How?**

Vijay Vignesh



**Top ten ways to tackle
overfitting models**

Shivani Shimpi



**Supervise.ly [Image
annotation and data
management]**

ANKUR DHAKAR



**Malaria Detection using
Deep Learning:**

Vishnu Dutt Pathak



Learn more.

Medium is an open platform where 170 million readers come to find insightful and dynamic thinking. Here, expert and undiscovered voices alike dive into the heart of any topic and bring new ideas to the surface. [Learn more](#)

Make Medium yours.

Follow the writers, publications, and topics that matter to you, and you'll see them on your homepage and in your inbox. [Explore](#)

Write a story on Medium.

If you have a story to tell, knowledge to share, or a perspective to offer — welcome home. It's easy and free to post your thinking on any topic. [Start a blog](#)

